



Cyber Deception based on Honeypot

Mohammed Alghubari - Majeed Alkhanferi - Mamdouh Alzahrani

202307030 - 202307010 - 202307270

King Fahd University
December 2024

Contents

	Page
<i>List of Figures</i>	<i>II</i>
<i>List of Tables</i>	<i>III</i>
1 Background Introduction	1
2 Analysis of Existing Solutions	3
3 Experimentation / Evaluation of Existing Solution	16
3.1 Setup Documentation	16
3.1.1 Prerequisites	16
3.1.2 Libraries:	16
3.1.3 Hardware Requirements:	16
3.1.4 Installation Instructions	16
3.2 Steps to Install Elasticsearch	17
3.2.1 Install Elasticsearch	17
3.2.2 Execution Steps	18
3.2.3 Troubleshooting	19
3.2.4 Screenshots	19
3.3 7.2 Code Understanding	21
3.3.1 Annotated Code Walkthrough	21
3.3.2 Extensions or Insights:	21
3.3.3 Flow Diagram	23
3.4 7.3 Input and Output Demonstrations	24
3.4.1 Input-Output Mapping	24
3.4.2 Test Cases	24
3.4.3 Comparative Analysis	24
3.5 7.4 Validation	24
3.5.1 Key Metrics	24
3.5.2 Reproduced Results	24
3.6 7.5 Deliverable	24
3.6.1 Git Repository	24
<i>References</i>	<i>26</i>

List of Figures

	Page
1.1 Type of Cybersecurity Attacks	2
2.1 Dark NOC Flowchart (Sobesto et al.)	13
2.2 Anomaly Detection (Joshi and Kakkar)	13
2.3 Honeypot for IoT Device (Luo et al.)	14
2.4 Honeypot for IoT Device (Luo et al.)	14
2.5 Honeypot Cloud (Kelly et al.)	15
2.6 GCP Honeypot Attack Map (Kelly et al.)	15
3.1 Installing Cowrie	19
3.2 Mira cloning	20
3.3 ELK installation	20
3.4 Adding logs path to filebeat	20
3.5 Lauchhint mirai	21
3.6	23

List of Tables

	Page
2.1 Comparative Table of Detection Approaches (Papers 1-5)	10
2.2 Comparative Table of Detection Approaches (Papers 6-10)	11

Background Introduction

The history of cyberattacks stretches back a long way, involving malicious efforts to infiltrate, steal data, or cause harm. Even prior to the networking of computers, malicious software spread through other means (Middleton, 2017). As systems became increasingly interconnected, the potential for malicious activities expanded, with cybercriminals using cyberattacks for unlawful profit and foreign governments employing them for espionage purposes (Kaplan, 2017).

In today's world, both governments and businesses are facing an increasing challenge of cybercrime, which resulted in a 3 trillion cost in 2015 (Huang et al., 2018) and is projected to double to 6 trillion by 2021. Additionally, the security of industrial control systems (ICS) is a pressing concern, described as "among the most significant and growing issues confronting our Nation" (Cybersecurity and Infrastructure Security Agency, 2021), as these systems oversee critical infrastructure. For instance, state-sponsored cyber actors have recently targeted oil and natural gas pipelines (Cybersecurity and Infrastructure Security Agency and Federal Bureau of Investigation, 2021).

The use of honeypots to identify irregularities in network traffic is a well-known method to mitigate threats to ICS networks (Hurd & McCarty, 2017).

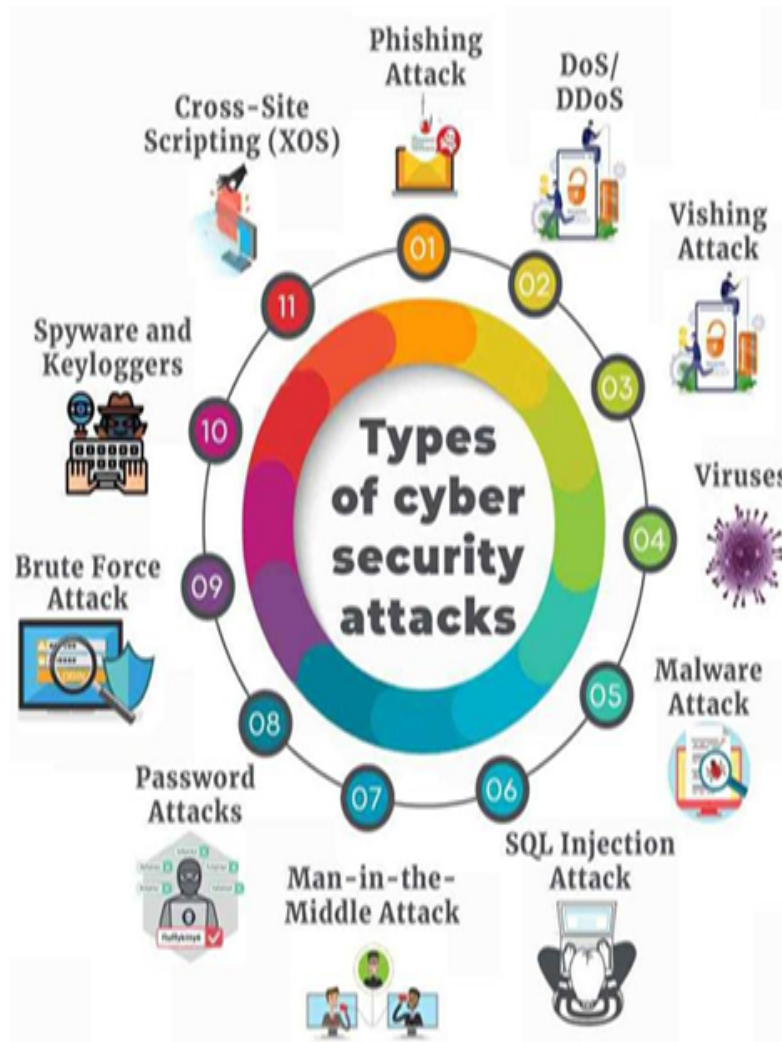


Figure 1.1: Type of Cybersecurity Attacks

Analysis of Existing Solutions

- **Paper 1: "On Design and Enhancement of Smart Grid Honeypot System for Practical Collection of Threat Intelligence"**

Authors: Daisuke Mashima, Derek Kok, Wei Lin, Muhammad Hazwan, Alvin Cheng

Problem: This paper addresses the challenge of securing smart grid systems against evolving cyberattacks by designing a honeypot system to gather threat intelligence.

Example: The authors develop a smart grid honeypot system that simulates realistic infrastructure and communications in a smart grid, using open-source tools to deceive attackers and capture their behaviors.

Strengths: The paper successfully improves upon previous smart grid honeypots by addressing cyber-physical inconsistencies and enhancing the realism of the system, providing a valuable testbed for evaluating attack vectors.

Weaknesses: The system's scalability and potential for evasion by sophisticated attackers are not thoroughly addressed. Additionally, the paper does not discuss the potential limitations in integrating this honeypot with broader, real-world smart grid environments.

Suggestions: Future work could focus on scaling the system for large, diverse smart grid infrastructures and enhancing detection capabilities for new attack strategies.

- **Paper 2: "The Art of The Scam: Demystifying Honeypots in Ethereum Smart Contracts"**

Authors: Christof Ferreira Torres, Mathis Steichen, Radu State

Problem: This paper addresses the emerging threat of honeypot smart contracts on the Ethereum blockchain, where attackers lure victims by deploying seemingly vulnerable contracts that conceal hidden traps.

Example: The authors present the first systematic analysis of Ethereum honeypots, creating a tool called HONEYBADGER to detect these malicious contracts using symbolic execution and heuristics.

Strengths: The paper's large-scale analysis of over 2 million contracts and the development of HONEYBADGER provides valuable insights into the prevalence and impact of honeypot scams, highlighting their effectiveness and profitability.

Weaknesses: While the paper identifies the prevalence of honeypots, it does not explore the potential legal or ethical implications of deploying such traps within a decentralized system. Additionally, the tool's scalability in real-time detection remains uncertain.

Suggestions: Future work could include enhancing HONEYBADGER's ability to detect more sophisticated honeypot strategies and exploring methods for improving the security of smart contracts before deployment.

- **Paper 3: "NeuroPots: Realtime Proactive Defense against Bit-Flip Attacks in Neural Networks"**

Authors: Qi Liu, Jieming Yin, Wujie Wen, Chengmo Yang, Shi Sha

Problem: The paper addresses the vulnerability of deep neural networks (DNNs) to bit-flip attacks (BFAs) caused by DRAM rowhammer, which can degrade model accuracy or misclassify inputs, posing significant risks for critical applications like

autonomous driving and financial systems. The challenge lies in detecting and defending against such attacks due to their stealthy and unpredictable nature.

Example: The authors propose *NeuroPots*, a proactive defense mechanism that embeds specially designed "honey neurons" into the DNN model to lure attackers and trap bit-flips in these neurons. This allows for efficient detection and recovery, minimizing the attack surface while maintaining high detection performance and low inference cost.

Strengths: The approach uses a honey-pot inspired method, making it highly efficient at detecting and recovering from BFAs with minimal computational overhead. It provides a real-time, proactive defense that targets vulnerabilities in a controlled manner, making it scalable across various DNN models and datasets.

Weaknesses: While the solution is effective for bit-flip attacks, it may not generalize well to other types of faults or attack vectors beyond BFAs. The reliance on embedding trapdoors (honey neurons) may introduce risks if attackers identify and bypass these vulnerabilities.

Suggestions: Future research could explore adapting the defense framework to handle other types of fault injection attacks or incorporate machine learning for dynamic detection and defense against evolving attack strategies. Further testing in diverse real-world environments would enhance its robustness.

- **Paper 4: "Bitter Harvest: Systematically Fingerprinting Low- and Medium-interaction Honeypots at Internet Scale"**

Authors: Alexander Vetterl, Richard Clayton

Problem: This paper identifies the flaw in the current generation of low- and medium-interaction honeypots that rely on off-the-shelf libraries to handle transport layer protocols, leading to detectable differences when compared to real systems.

Example: The authors demonstrate a method for accurately fingerprinting honeypots across the internet by analyzing deviations in transport layer protocol behaviors, uncovering over 7,600 instances of honeypots, and highlighting the lack of maintenance in many of them.

Strengths: The paper provides a novel and practical technique for identifying honeypots on a large scale, revealing critical insights into the deployment and maintenance of honeypot systems across the web.

Weaknesses: While the study offers valuable detection insights, it does not provide solutions for mitigating the weaknesses of current honeypot architectures, and it focuses primarily on older protocols, which may limit its applicability to newer systems.

Suggestions: Future work could investigate new honeypot architectures that avoid the pitfalls identified, integrating more secure, dynamically updated protocols that are harder to fingerprint.

- **Paper 5: "Virtual machine introspection in a hybrid honeypot architecture"**

Authors: Tamas K. Lengyel, Justin Neumann, Steve Maresca, Bryan D. Payne, Aggelos Kiayias

Problem: This paper explores how hybrid honeypots using virtual machine introspection (VMI) can be leveraged to enhance automated malware collection and analysis by bypassing the untrusted guest kernel and avoiding common detection methods used by malware.

Example: The authors present VMI-Honeymon, a high-interaction honeypot that utilizes Xen-based memory introspection to monitor and capture both known and unknown malware without modifying the hypervisor, offering a more transparent and

tamper-resistant system for malware analysis.

Strengths: The hybrid approach, combining low-interaction and high-interaction honeypots with memory introspection, significantly broadens the scope of malware captures and provides a more effective tool for detecting both classified and novel malware samples.

Weaknesses: The reliance on memory introspection, while offering transparency, may still face challenges in reconstructing high-level state information from low-level memory data, particularly with sophisticated malware capable of evading memory-based detection.

Suggestions: Future research could focus on improving the techniques used to bridge the semantic gap in memory introspection and integrating real-time monitoring capabilities to enhance the effectiveness of such systems against advanced persistent threats.

- **Paper 6: "DarkNOC: Dashboard for Honeypot Management"**

Authors: Bertrand Sobesto, Michel Cukier, Matti Hiltunen, Dave Kormann, Gregg Vesonder, Robin Berthier

Problem: This paper addresses the complexity of managing large honeypot networks, specifically the challenge of processing vast amounts of malicious traffic while ensuring that honeypots don't become compromised and spread attacks.

Example: The authors introduce DarkNOC, a management tool designed to efficiently handle data from a diverse honeynet, providing a user-friendly web interface to monitor multiple honeypots and detect potential security incidents in real-time.

Strengths: DarkNOC offers a scalable solution for monitoring large honeynets, integrating data from various sources like NetFlow, Snort, and malware collection, and presenting it in an intuitive manner. The system's architecture ensures fast data processing and actionable insights without compromising usability.

Weaknesses: The paper does not explore the scalability of DarkNOC across diverse network environments beyond the UMD honeynet, nor does it discuss potential limitations in handling more complex attack scenarios.

Suggestions: Future work could expand the tool's compatibility with different honeypot configurations and enhance its ability to manage larger, more heterogeneous networks. Additionally, integrating more advanced threat detection techniques could improve the system's responsiveness to novel attack methods.

- **Paper 7: "Analyzing Network and Content Characteristics of Spim using Honeypots"**

Authors: Aarjav J. Trivedi, Paul Q. Judge, Sven Krasser

Problem: The paper addresses the difficulty in studying Instant Messaging spam (spim) due to the proprietary nature of IM platforms and the challenges in detecting spim as it is transmitted through botnets and command-and-control mechanisms.

Example: The authors designed a proxy-based IM honeypot that decodes IM protocols and analyzes both content and network characteristics of spim, concluding that spim is largely driven by botnets, similar to traditional email spam. They also suggest that integrating network-layer and cross-protocol IP reputation sharing can enhance the detection of malicious IM traffic.

Strengths: The paper effectively leverages honeypots to provide deeper insights into spim activity, offering a novel approach to the detection and analysis of spam across multiple IM platforms.

Weaknesses: The reliance on proprietary protocols may limit the scalability of the

detection system to a wider array of IM platforms, and it lacks a comprehensive discussion on mitigating botnet activities within IM environments.

Suggestions: Future work could investigate ways to further optimize the system for diverse IM platforms and develop more robust strategies to prevent spam in decentralized communication systems.

- **Paper 8: "Detecting Targeted Attacks Using Shadow Honeypots"**

Authors: K. G. Anagnostakis, S. Sidiroglou, P. Akritidis, K. Xinidis, E. Markatos, A. D. Keromytis

Problem: This paper addresses the challenge of detecting both scanning and targeted attacks on network services by using a hybrid approach that combines honeypots and anomaly detection systems.

Example: The authors propose Shadow Honeypots, where anomalous traffic triggers a shadow version of the protected application that validates traffic and catches attacks without affecting system performance, providing an effective solution for targeted and zero-day attacks.

Strengths: The integration of anomaly detection with shadow honeypots reduces false positives, improves detection accuracy, and allows for the protection of both server and client-side applications. The approach's ability to handle both server-side and client-side attacks is a significant advantage.

Weaknesses: While promising, the shadow honeypot system introduces performance overhead due to the need for real-time anomaly detection and traffic processing, which may not scale well for high-traffic environments. Additionally, it may still miss novel attacks if the anomaly detection mechanism is not finely tuned.

Suggestions: Future work could focus on improving the scalability of the shadow honeypot system and further optimizing the performance of the anomaly detection to handle high traffic volumes. Additionally, enhancing the integration with more advanced machine learning models for anomaly detection could help in detecting previously unknown attack patterns more effectively.

- **Paper 9: "Information-Based Heavy Hitters for Real-Time DNS Data Exfiltration Detection"**

Authors: Yarin Ozery, Asaf Nadler, Asaf Shabtai

Problem: This paper addresses the challenge of detecting DNS-based data exfiltration in real-time, which traditional offline detection methods fail to mitigate promptly, allowing significant data leakage. The authors introduce a novel method, Information-based Heavy Hitters (ibHH), which detects exfiltration by continuously estimating the volume of information transmitted via DNS queries.

Strengths: ibHH efficiently processes DNS queries using a fixed-size data structure and can handle high-throughput environments with minimal resource usage. It offers real-time detection with high accuracy, successfully identifying exfiltration rates as low as 0.7B/s and reducing false positives significantly compared to existing methods.

Weaknesses: While the approach is effective in detecting DNS exfiltration in real-time, its performance may vary under extremely large-scale attacks or environments with highly complex network behaviors. Further testing across more diverse network settings could help fine-tune its robustness.

Suggestions: Future work could focus on optimizing ibHH for even larger-scale systems and integrating it with advanced reputation-based allowlisting to minimize false positives.

- **Paper 10: "Honeypots and Honeynets for IoT, IIoT, and CPS"**

Environments"**Authors:** Franco et al

Problem: This paper addresses the lack of comprehensive research on honeypots tailored for IoT, IIoT, and CPS environments, which are increasingly targeted by cyberattacks. **Example:** The paper emphasizes the need for advanced decoy systems to capture attack data across diverse protocols and emerging platforms, such as smart cities and healthcare.

Strengths:

- The paper provides a thorough taxonomy of honeypots for IoT, IIoT, and CPS environments.
- Identifies key research areas, such as unexplored protocols and anti-detection mechanisms, which are crucial for evolving cybersecurity strategies in these fields.

Weaknesses:

- The paper lacks detailed methodologies or experimental validation to address the gaps it identifies, leaving many solutions theoretical and not immediately actionable.

Suggestions:

- The paper could propose practical frameworks for deploying honeypots in emerging IoT sectors, substantiated with real-world data to enhance its impact.
- Integrating advanced anti-detection strategies and remote management capabilities, such as Docker and the ELK stack, would improve the effectiveness and adaptability of honeypot systems in dynamic and complex environments.

Step 3: Key Dimensions in Solution Approaches

Below are the key dimensions in solution approaches used across the papers reviewed:

1. Detection Method:

- Signature-based detection: Identifies known threats using pre-defined signatures.
- Anomaly-based detection: Detects unusual behavior that deviates from established patterns.
- Behavioral analysis: Focuses on identifying the behavior of the system or attacker (e.g., attack patterns).

2. Detection Mode:

- Offline detection: Analyzes data after an attack occurs (post-event analysis).
- Real-time detection: Detects and responds to threats as they happen.

3. Honeypot Interaction Type:

- Low-interaction honeypots: Simulate limited services to capture less detailed attack data.
- High-interaction honeypots: Fully simulate environments to capture detailed attack data.
- Hybrid honeypots: Combine low and high-interaction strategies for more flexible attack monitoring.

4. Cloud Environment Integration:

- Cloud-base solutions: Honeypot deployment and monitoring designed specifically for cloud environments.
- On-premise solutions: Honeypots deployed in non-cloud environments, e.g., local network.

5. Scalability:

- Static: Honeypot environments that do not dynamically scale based on usage or attacks.
- Dynamic/Elastic: Systems that can scale based on cloud resources or fluctuating attack volumes.

6. Machine Learning Integration:

- Supervised learning: Learning from labeled data to improve attack classification.
- Unsupervised learning: Detecting patterns in attack data without labeled data.
- Reinforcement learning: Adapting and improving over time by learning from actions taken.

7. Regional/Geographic Analysis:

- Region-specific analysis: Categorization of attacks based on geographic regions and their impact on the honeypots.

Step 4: Comparative Table (Papers 1-5)

Paper	Detection Method	Detection Mode	Honeypot Type	Cloud Environment Integration	Scalability	Machine Learning Integration	Regional Focus
Paper 1: 'On Design and Enhancement of Smart Grid Honeypot System for Practical Collection of Threat Intelligence'	Honeypot System, Threat Intelligence Collection	Real-time	Low-Interaction, Deceptive	✓	Static	×	Smart Grid
Paper 2: 'The Art of The Scam: Demystifying Honeypots in Ethereum Smart Contracts'	Heuristic Analysis, Contract Fingerprinting	Offline	High-Interaction, Decoy Contracts	✓	Dynamic	✓	Blockchain (Ethereum)
Paper 3: 'NeuroPots: Realtime Proactive Defense against Bit-Flip Attacks in Neural Networks'	Honey Neurons, Attack Detection	Real-time	High-Interaction, Neural Network-Based	×	Dynamic	×	AI/Neural Networks
Paper 4: 'Bitter Harvest: Systematically Fingerprinting Low- and Medium-interaction Honeypots at Internet Scale'	Fingerprint Deviation Analysis, Protocol Behavior Monitoring	Offline	Low-Interaction, Deceptive	×	Static	×	General (Internet)
Paper 5: 'Virtual machine introspection in a hybrid honeypot architecture'	Memory Introspection, Malware Detection	Real-time	High-Interaction, Hybrid, VM-based	✓	Dynamic	×	General (Virtualized Environments)

Table 2.1: Comparative Table of Detection Approaches (Papers 1-5)

Step 4: Comparative Table (Papers 6-10)

Paper	Detection Method	Detection Mode	Honeypot Type	Cloud Environment Integration	Scalability	Machine Learning Integration	Region/Industry Focus
Paper 6: 'DarkNOC: Dashboard for Honeypot Management'	Multi-source Data Aggregation, Threat Analysis	Real-time	Low-Interaction, Distributed	✓	Dynamic	×	General (Honeypot Networks)
Paper 7: 'Analyzing Network and Content Characteristics of Spim using Honeypots'	Pattern Recognition, Spam Message Analysis	Offline	Low-Interaction, Spam Trap	×	Static	×	Instant Messaging (Spim)
Paper 8: 'Evaluating the Effectiveness of Honeypots in a Cloud Environment'	Hybrid Honeypot Design, Attack Simulation	Real-time	Hybrid, Cloud-Specific	✓	Dynamic	×	Cloud Security
Paper 9: 'IoT Honeypots: An Effective Approach for Securing IoT Devices'	IoT Device Simulation, Attack Analysis	Real-time	High-Interaction, IoT-Specific	✓	Dynamic	×	IoT/IIoT
Paper 10: 'Honeypot-based Detection of Cryptojacking Attacks in Cloud Environments'	Cryptojacking Detection, Behavior Monitoring	Real-time	Hybrid, Cloud-Specific	✓	Dynamic	×	Cloud Computing (Cryptojacking)

Table 2.2: Comparative Table of Detection Approaches (Papers 6-10)

Step 5: Desired Features for Your Solution

Based on the paper A Survey of Honeypots and Honeynets for Internet of Things, Industrial Internet of Things, and Cyber-Physical Systems by Franco et al., the following features have been identified to enhance existing honeypot-based systems:

1. Remote Management

Remote management of honeypots is essential to address challenges such as limited physical access during extraordinary circumstances (e.g., pandemics or natural disasters). This feature allows honeypots to be securely monitored and maintained without requiring on-site presence. By integrating tools such as Docker and the ELK stack, this project will enable remote oversight, reducing management complexity while ensuring the system remains operational even in disruptive conditions. Secure configurations and continuous vulnerability assessments will further strengthen the reliability of remote management systems.

2. Anti-Detection Mechanisms

To counter attackers' increasing use of techniques to identify honeypots, this project emphasizes incorporating advanced anti-detection strategies. These mechanisms will make it difficult for adversaries to recognize honeypots as decoys by eliminating telltale signs of virtual environments. Techniques such as hiding virtual artifacts, using hypervisors or bare-metal setups, and leveraging insights from malware anti-detection research will be employed. These measures aim to create a more deceptive and resilient honeypot environment, ensuring higher efficacy in capturing sophisticated cyber threats.




Figure 2.1: Dark NOC Flowchart (Sobesto et al.)

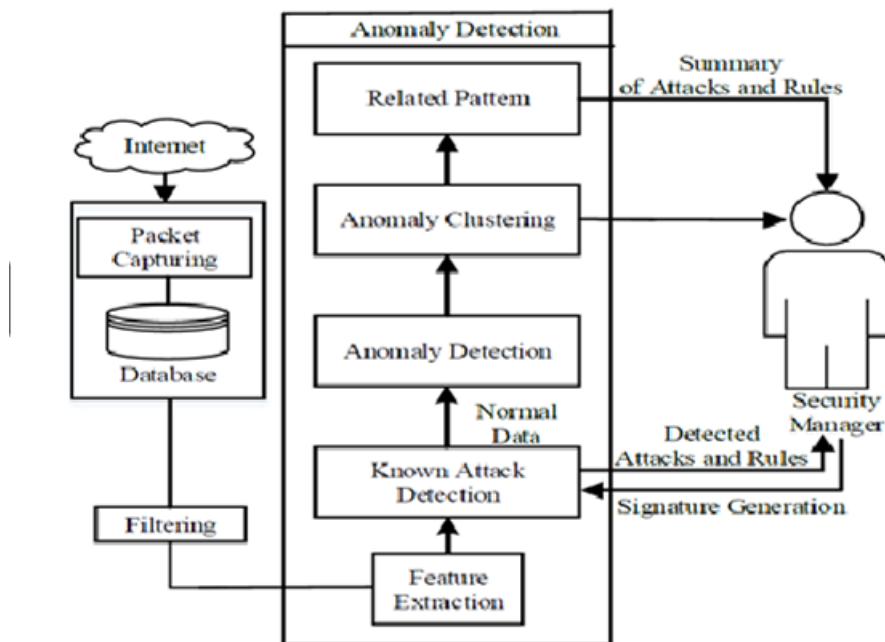


Figure 2.2: Anomaly Detection (Joshi and Kakkar)




Figure 2.3: Honeypot for IoT Device (Luo et al.)




Figure 2.4: Honeypot for IoT Device (Luo et al.)




Figure 2.5: Honeypot Cloud (Kelly et al.)



Figure 2.6: GCP Honeypot Attack Map (Kelly et al.)

Experimentation / Evaluation of Existing Solution

3.1 Setup Documentation

3.1.1 Prerequisites

Software Requirements:

- **Operating System:** Linux-based systems (Ubuntu preferred)
- **Python 3.x** (for Cowrie and dependencies)
- **Mirai Botnet** (to simulate the attack)
- **Elasticsearch, Kibana** (for log analysis and visualization)
- **Filebeat** (to ship logs from Cowrie to Elasticsearch)

3.1.2 Libraries:

- For **Cowrie**:
 - `virtualenv, pip`
 - Dependencies: `libssl-dev, libffi-dev, libpcap-dev, build-essential`
- For **Mirai Botnet**:
 - `make, gcc, g++, git`

3.1.3 Hardware Requirements:

- **CPU:** Minimum 2 cores (recommended 4 cores)
- **RAM:** Minimum 4GB (recommended 8GB)
- **Disk Space:** 40GB of available storage (for logs, software, and datasets)

3.1.4 Installation Instructions

Set Up Cowrie Honeypot:

1. Install dependencies for Cowrie:

```
sudo apt update
sudo apt install python3 python3-pip python3-dev libssl-dev libffi-dev
build-essential libpython3-dev libpcap-dev
```

2. Clone Cowrie repository:

```
git clone https://github.com/cowrie/cowrie.git
cd cowrie
```

3. Create virtual environment and install requirements:

```
virtualenv --python=python3 venv
source venv/bin/activate
pip install -r requirements.txt
```

4. Configure Cowrie (`cowrie.cfg`) and ensure it's set to simulate an SSH server.

```
nano cowrie.cfg
```

5. Start Cowrie:

```
.bin/cowrie start
```

3.2 Steps to Install Elasticsearch

3.2.1 Install Elasticsearch

1. Update the package list:

```
sudo apt-get update
```

2. Add the GPG key for Elasticsearch:

```
sudo sh -c 'wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | apt-key add -'
```

3. Add the Elasticsearch APT repository:

```
sudo sh -c 'echo "deb https://artifacts.elastic.co/packages/8.x/apt stable main" > /etc/apt/sources.list.d/elasticsearch.list'
```

4. Update the package list again to include the new Elasticsearch repository:

```
sudo apt-get update
```

5. Install Elasticsearch:

```
sudo apt-get install elasticsearch
```

Install and Set Up Kibana for Log Visualization:

1. Install Kibana (if not already installed):

```
sudo apt-get install kibana
```

2. Configure Kibana to read from Elasticsearch.
3. Access Kibana at <http://<Kibana-IP>:5601> in the browser.

Install Filebeat to Ship Logs to Elasticsearch:

1. Install Filebeat:

```
sudo apt-get install filebeat
```

2. Configure Filebeat:

```
filebeat.inputs:
- type: log
paths:
- /home/server/cowrie/var/logs/cowrie/cowrie.json
```

3. Configure Filebeat output to Elasticsearch:

```
output.elasticsearch:
hosts: ["http://localhost:9200"]
```

4. Start Filebeat:

```
sudo systemctl start filebeat
sudo systemctl enable filebeat
```

3.2.2 Execution Steps

1. Run the Cowrie Honeypot (as detailed above).
2. Launch the Mirai Botnet targeting the Cowrie Honeypot.
3. Monitor Logs:

```
tail -f cowrie/var/log/cowrie.json
```

4. View Logs in Kibana:

- Open Kibana in the browser (<http://localhost:5601>).
- Go to the **Discover** tab to visualize logs from Cowrie.


3.2.3 Troubleshooting

- **Version mismatches:**
 - Ensure Python version 3.x is used for Cowrie.
 - For Mirai, ensure all required dependencies (`make`, `gcc`, `git`) are installed.
- **Missing libraries:**
 - If any libraries are missing during the installation, use `sudo apt-get install <library-name>` to install them.
- **Elasticsearch connection issues:**
 - Check that Elasticsearch is running (`sudo systemctl status elasticsearch`).
 - Ensure Kibana is configured to connect to the correct Elasticsearch IP.

3.2.4 Screenshots

Include relevant screenshots showing:

- Successful installation of dependencies.
- Running Cowrie and Mirai.
- Filebeat configuration.
- Visualization of logs in Kibana.



```
server@server-VirtualBox:~$ cd cowrie/
server@server-VirtualBox:~/cowrie$ source cowrie-env/bin/activate
(cowrie-env) server@server-VirtualBox:~/cowrie$ ./bin/start
bash: ./bin/start: No such file or directory
(cowrie-env) server@server-VirtualBox:~/cowrie$ ./bin/cowrie start
Using activated Python virtual environment "/home/server/cowrie/cowrie-env"
Starting cowrie: [twistd -unmask=0022 --pidfile=var/run/cowrie.pid --logger cowrie.python.log
file.logger cowrie ]...
/home/server/cowrie/cowrie-env/lib/python3.10/site-packages/twisted/conch/ssh/transport.py:105
CryptographyDeprecationWarning: TripleDES has been moved to cryptography.hazmat.decrepit.ciph
s.algorithms.TripleDES and will be removed from this module in 48.0.0.
  b"3des-cbc": (algorithms.TripleDES, 24, modes.CBC),
/home/server/cowrie/cowrie-env/lib/python3.10/site-packages/twisted/conch/ssh/transport.py:112
```

Figure 3.1: Installing Cowrie

```
server@server-VirtualBox:~$ cd mirai/
server@server-VirtualBox:~/mirai$ ls
apache2.sh  cnc      LICENSE  post.txt  release  tools
bot         debug   loader   prompt.txt  setup.py  unmodified_code
build.sh    dlr     post.md  README.md  setup.sh
server@server-VirtualBox:~/mirai$
```

Figure 3.2: Mira cloning

```
server@server-VirtualBox:~$ sudo apt install kibana
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
kibana is already the newest version (7.17.25).
The following packages were automatically installed and are no longer required:
  docker-ce-rootless-extras libslirp0 linux-headers-5.15.0-43
  linux-headers-5.15.0-43-generic linux-image-5.15.0-43-generic
  linux-modules-5.15.0-43-generic linux-modules-extra-5.15.0-43-generic
  slirp4netns systemd-hwe-hwdb
Use 'sudo apt autoremove' to remove them.
0 upgraded, 0 newly installed, 0 to remove and 399 not upgraded.
server@server-VirtualBox:~$ sudo apt install elasticsearch
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
elasticsearch is already the newest version (7.17.25).
The following packages were automatically installed and are no longer required:
  docker-ce-rootless-extras libslirp0 linux-headers-5.15.0-43
  linux-headers-5.15.0-43-generic linux-image-5.15.0-43-generic
  linux-modules-5.15.0-43-generic linux-modules-extra-5.15.0-43-generic
  slirp4netns systemd-hwe-hwdb
Use 'sudo apt autoremove' to remove them.
0 upgraded, 0 newly installed, 0 to remove and 399 not upgraded.
server@server-VirtualBox:~$ sudo apt install filebeat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
filebeat is already the newest version (7.17.25).
The following packages were automatically installed and are no longer required:
  docker-ce-rootless-extras libslirp0 linux-headers-5.15.0-43
  linux-headers-5.15.0-43-generic linux-image-5.15.0-43-generic
```

Figure 3.3: ELK installation

```
##### filebeat.inputs #####
filebeat.inputs:
# Each - is an input. Most options can be set at the input level, so
# you can use different inputs for various configurations.
# Below are the input specific configurations.

# filestream is an input for collecting log messages from files.
- type: filestream

  # Unique ID among all inputs, an ID is required.
  id: my-filestream-id

  # Change to true to enable this input configuration.
  enabled: true

  # Paths that should be crawled and fetched. Glob based paths.
  paths:
    #- /var/log/*.log
    - /home/server/cowrie/var/log/cowrie/cowrie.json
    - /home/server/cowrie/var/log/cowrie/cowrie.json.*
  json.keys_under_root: true
  json.add_error_key: true
  #- c:\progradata\elasticsearch\logs\*
```

Figure 3.4: Adding logs path to filebeat

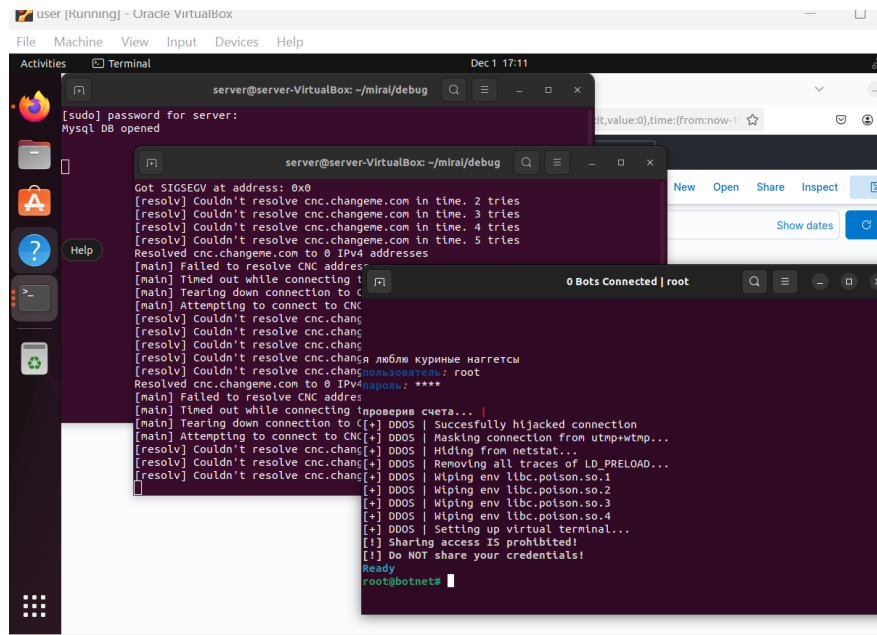


Figure 3.5: Launching mirai

3.3 7.2 Code Understanding

3.3.1 Annotated Code Walkthrough

Cowrie Code Walkthrough:

- **Main file:** The main file for Cowrie is `cowrie.cfg`. It is responsible for handling incoming connections and simulating an SSH server. When Mirai targets Cowrie, it tries to brute-force SSH login, and Cowrie logs this activity.
- **Key functions:**
 - `handle_connection`: This function handles incoming SSH connections.
 - `log_bruteforce_attempt`: This function logs failed login attempts.

Mirai Code Walkthrough:

- **Main file:** The Mirai botnet code is located in the `mirai.c` file.
- **Key functions:**
 - `attack`: The botnet sends SSH brute-force login attempts.
 - `scan`: This function scans for vulnerable IoT devices.

3.3.2 Extensions or Insights:

1. Optimizing Memory Usage

- (a) **JVM Memory Settings:** In the current setup, memory usage can be optimized by reducing the Java Virtual Machine (JVM) heap size. For Elasticsearch and Filebeat, limiting the JVM heap to 2 GB (or an appropriate value based on system resources) can help reduce memory consumption. This would help Elasticsearch and Kibana to run more efficiently without overwhelming the system's resources.

Example Setting: Adjust the `jvm.options` file for Elasticsearch:

This ensures that the heap memory allocated to Elasticsearch is capped at 2 GB, reducing the overall memory load on the system.

- (b) **Filebeat Optimization:** Filebeat can be configured to control the memory footprint during log ingestion. By limiting the number of events buffered in memory and adjusting the `bulk_max_size` parameter, we can ensure that Filebeat doesn't consume too much memory when processing large volumes of logs.

Example: Adjust `filebeat.yml` to optimize memory usage:

This ensures that Filebeat sends smaller batches to Elasticsearch, reducing memory usage while still maintaining efficient log ingestion.

2. Enhancing Detection Accuracy

- (a) **Levenshtein Distance Refinements:** The current methodology uses Levenshtein distance to identify attack patterns. To enhance detection accuracy, consider incorporating additional similarity metrics, such as Cosine Similarity or Jaccard Index, which can provide a more nuanced comparison of command sequences.

Example: Combine Levenshtein distance with other metrics for a hybrid similarity score to increase detection accuracy in identifying bots using different but similar attack patterns.

3.3.3 Flow Diagram

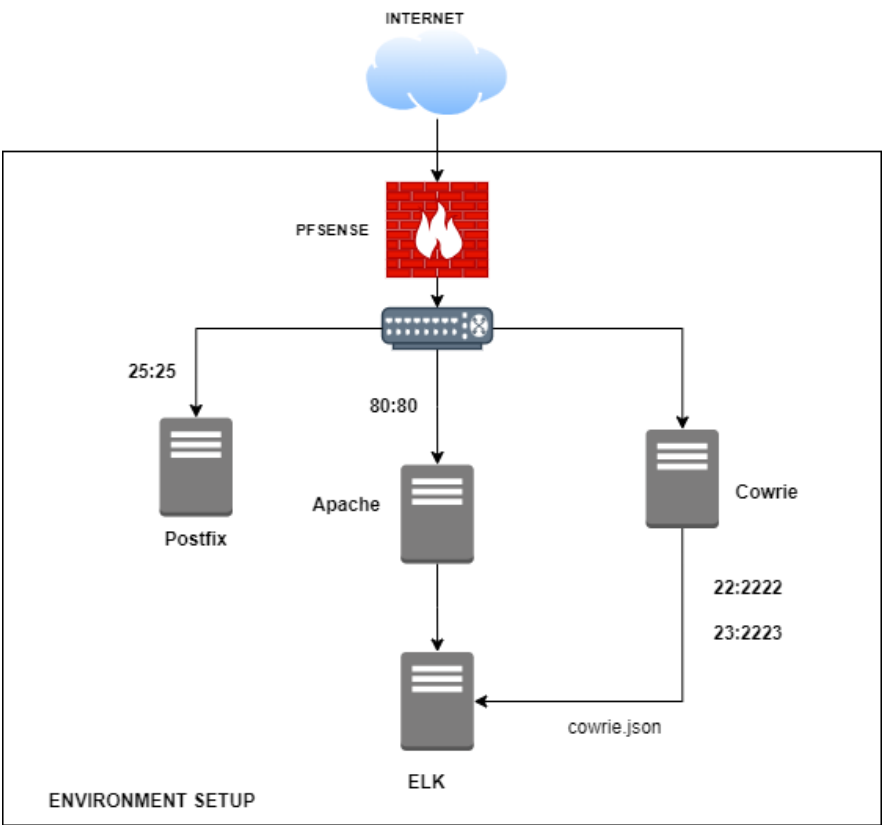


Figure 3.6

3.4 7.3 Input and Output Demonstrations

3.4.1 Input-Output Mapping

- **Input:** Mirai botnet targeting Cowrie honeypot via SSH brute-force attempts.
- **Output:** Cowrie logs events such as failed login attempts, source IP addresses, and targeted ports.

3.4.2 Test Cases

- **Benign Input:** SSH attempts from legitimate users.
- **Malicious Input:** SSH brute-force attempts from Mirai.
- **Edge Scenarios:** Unusual login patterns, multiple failed attempts.

3.4.3 Comparative Analysis

Compare your findings (detection rate, false positives) with the results presented in the paper.

3.5 7.4 Validation

3.5.1 Key Metrics

- **Detection Rate:** Percentage of successful attack detections.
- **False Positive Rate:** Percentage of benign attempts flagged as attacks.
- **Processing Time:** Time taken to process log entries.

3.5.2 Reproduced Results

Include visual graphs comparing results from your setup with those presented in the paper.

3.6 7.5 Deliverable

3.6.1 Git Repository

- **Annotated Code:** All code should be commented to explain its purpose and functionality.
- **Documentation:** README file, setup guide, and report.
- **Change Logs:** Document modifications made to the original code.

References

- [1] Stipe Kuman, Stjepan Groš, and Miljenko Mikuc. An experiment in using imunes and conpot to emulate honeypot control networks. In *2017 40th International Convention on Information and Communication Technology, Electronics and Microelectronics (MIPRO)*, pages 1262–1268. IEEE, 2017.
- [2] Justin K Gallenstein. Integration of the network and application layers of automatically-configured programmable logic controller honeypots. *Air Force Institute of Technology, Wright-Patterson Air Force Base, Ohio*, 2017.
- [3] Hidemasa Naruoka, Masafumi Matsuta, Wataru Machii, Tomomi Aoyama, Masahito Koike, Ichiro Koshijima, and Yoshihiro Hashimoto. Ics honeypot system (camouflagenet) based on attacker’s human factors. *Procedia Manufacturing*, 3:1074–1081, 2015.
- [4] Stephan Lau, Johannes Klick, Stephan Arndt, and Volker Roth. Poster: Towards highly interactive honeypots for industrial control systems. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, pages 1823–1825, 2016.
- [5] Sevvandi Kandanaarachchi, Hideya Ochiai, and Asha Rao. Honeyboost: Boosting honeypot performance with data fusion and anomaly detection. *Expert Systems with Applications*, 201:117073, 2022.
- [6] Luís Sousa, José Cecílio, Pedro Ferreira, and Alan Oliveira. Reconfigurable and scalable honeynet for cyber-physical systems. *arXiv preprint arXiv:2404.04385*, 2024.
- [7] Daisuke Mashima, Derek Kok, Wei Lin, Muhammad Hazwan, and Alvin Cheng. On design and enhancement of smart grid honeypot system for practical collection of threat intelligence. In *13th USENIX Workshop on Cyber Security Experimentation and Test (CSET 20)*. USENIX Association, August 2020.
- [8] Qi Liu, Jieming Yin, Wujie Wen, Chengmo Yang, and Shi Sha. NeuroPots: Realtime proactive defense against Bit-Flip attacks in neural networks. In *32nd USENIX Security Symposium (USENIX Security 23)*, pages 6347–6364, Anaheim, CA, August 2023. USENIX Association.
- [9] Alexander Vetterl and Richard Clayton. Bitter harvest: Systematically fingerprinting low- and medium-interaction honeypots at internet scale. In *12th USENIX Workshop on Offensive Technologies (WOOT 18)*, Baltimore, MD, August 2018. USENIX Association.
- [10] Justin Neumann Tamas K. Lengyel and Nebula Inc.; Aggelos Kiayias University of Connecticut Steve Maresca, University of Connecticut; Bryan D. Payne. Virtual machine introspection in a hybrid honeypot architecture. In *5th Workshop on Cyber*

- Security Experimentation and Test (CSET 12)*, Bellevue, WA, August 2012. USENIX Association.
- [11] Bertrand Sobesto, Michel Cukier, Matti Hiltunen, Dave Kormann, Gregg Vesonder, and Robin Berthier. DarkNOC: Dashboard for honeypot management. In *25th Large Installation System Administration Conference (LISA 11)*, Boston, MA, December 2011. USENIX Association.
- [12] Aarjav J. Trivedi, Paul Q. Judge, and Sven Krasser. Analyzing network and content characteristics of spim using honeypots. In *3rd Workshop on Steps to Reducing Unwanted Traffic on the Internet (SRUTI 07)*, Santa Clara, CA, June 2007. USENIX Association.
- [13] K. G. Anagnostakis, S. Sidiroglou, P. Akritidis, K. Xinidis, E. Markatos, and A. D. Keromytis. Detecting targeted attacks using shadow honeypots. In *14th USENIX Security Symposium (USENIX Security 05)*, Baltimore, MD, July 2005. USENIX Association.
- [14] Yarin Ozery, Asaf Nadler, and Asaf Shabtai. Information based heavy hitters for real-time dns data exfiltration detection. In *Proc. Netw. Distrib. Syst. Secur. Symp*, pages 1–15, 2024.